

Hands-on Introduction to TensorFlow 2.0



Josh Gordon (@random_forests)

Robert Crowe (@robert_crowe)

At ICSE Montreal



Getting started

You will need

- A laptop with an internet connection.
- There's nothing to install in advance.

Agenda

- **Beginner** examples in the first half.
- **Advanced** examples in the second.

Goal

- Write and run a few basic examples, know where to go to learn more.



Agenda

Beginner examples

- Installing TensorFlow 2.0 and using Colab
- Linear regression
- MNIST (for beginners); MNIST (for researchers)

Intermediate examples

- Serving a model with a REST API
- Serving a model in a browser

Advanced examples (as time remains)

- Deep Dream, Neural machine translation, Image Colorization



TensorFlow

An open source Deep Learning library

- Released by Google in 2015
- >1800 contributors worldwide

TensorFlow 2.0

- Easier to use
- For beginners and experts
- In alpha, final version coming in a few months



TensorFlow is basically

NumPy

- + GPU / TPU support
- + AutoDiff
- + Utilities to help you write neural networks (layers, optimizers)

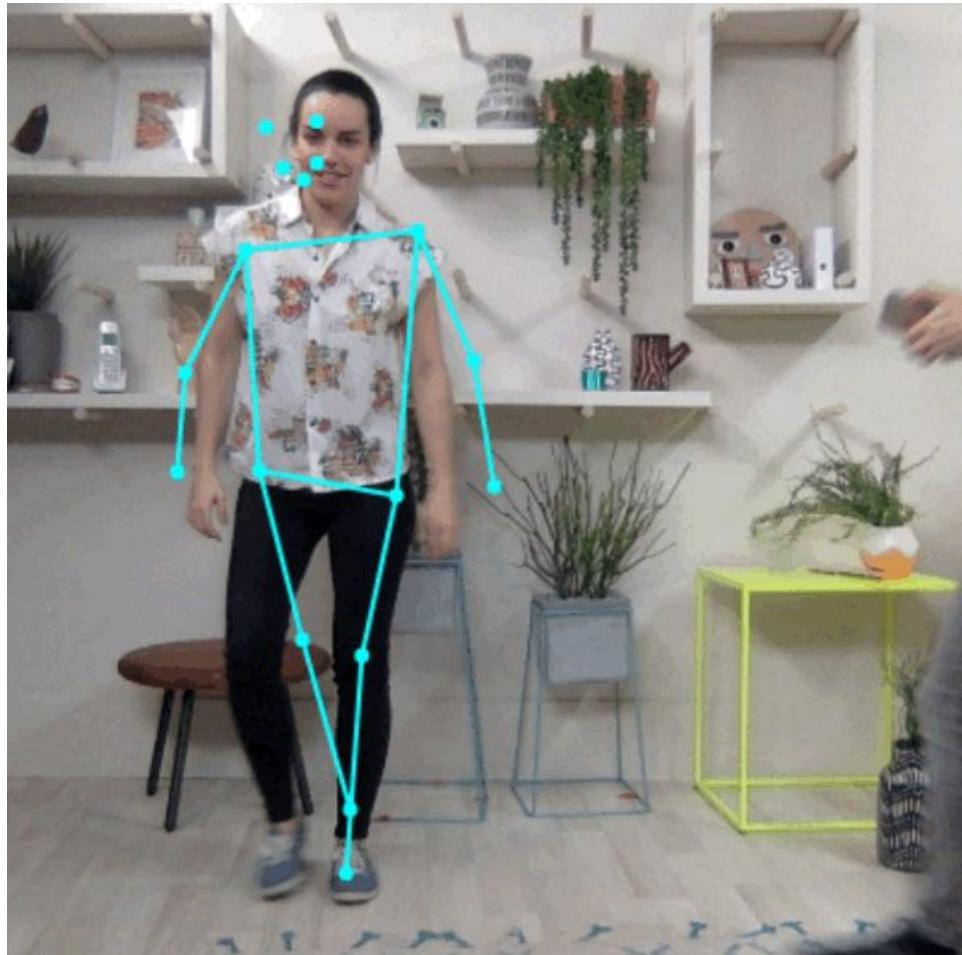
TensorFlow

- A C++ engine to accelerate code written in Python.
- Your program is represented as a **graphs** that can run on devices **without a Python interpreter (phones, browsers, etc)**



PoseNet

bit.ly/pose-net





The person is riding a surfboard in the waves.



Topics you'll learn about

Building models

- Sequential models
- Subclassed models

Under the hood

- AutoGraph and `tf.function`
- TF2 vs TF1

Training models

- Built-in vs custom training loops

Learning more

- Book recommendations

You can use TF 2.0 like NumPy

```
import tensorflow as tf # Assuming TF 2.0 is installed

a = tf.constant([[1, 2], [3, 4]])
b = tf.matmul(a, a)

print(b)
# tf.Tensor( [[ 7 10] [15 22]], shape=(2, 2), dtype=int32)

print(type(b.numpy()))
# <class 'numpy.ndarray'>
```



What's Deep Learning?

Representation learning



What is Deep Learning?

Quick demos with [TensorFlow Playground](#)

playground.tensorflow.org

DATA

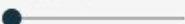
Which dataset do you want to use?



Ratio of training to test data: 50%



Noise: 0



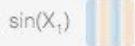
Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?



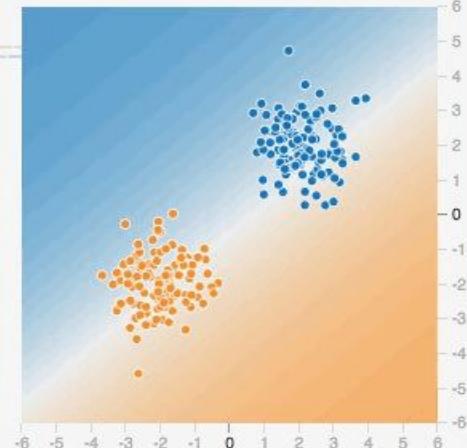
+

-

0 HIDDEN LAYERS

OUTPUT

Test loss 0.455
Training loss 0.455



Colors shows
data, neuron and
weight values.



Show test data

Discretize output

DATA

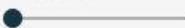
Which dataset do you want to use?



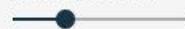
Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?

X_1



X_2



X_1^2



X_2^2



$X_1 X_2$



$\sin(X_1)$



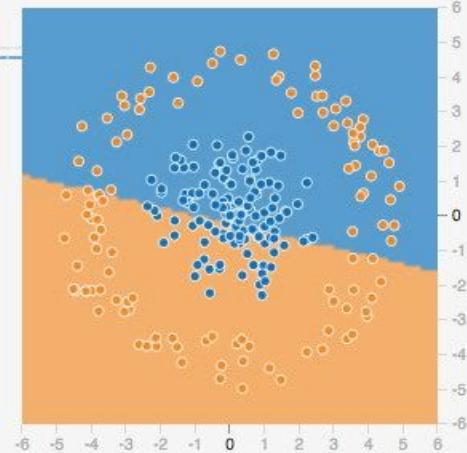
$\sin(X_2)$



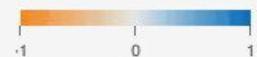
0 HIDDEN LAYERS

OUTPUT

Test loss 0.672
Training loss 0.627

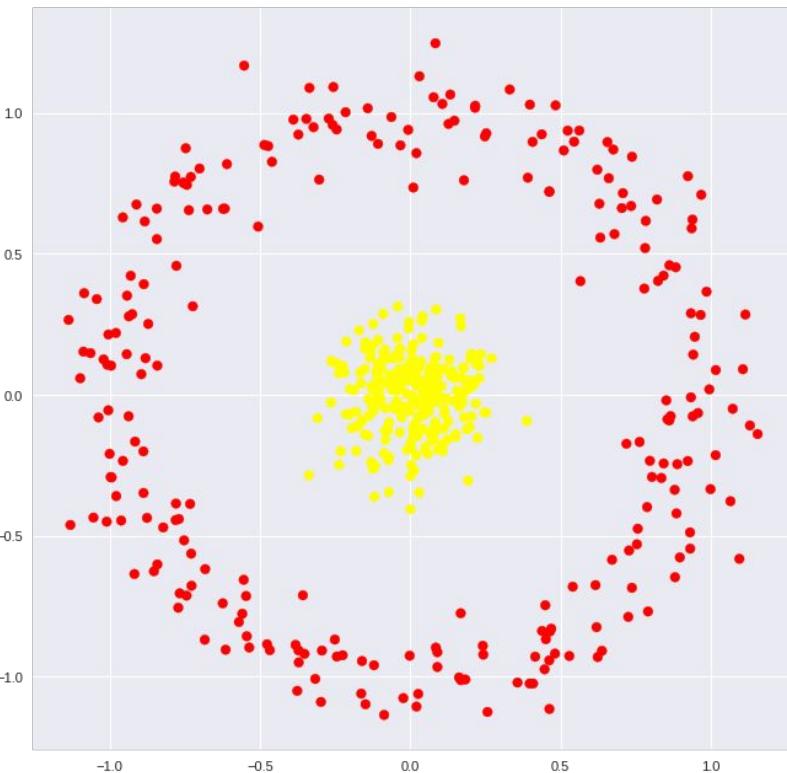


Colors shows data, neuron and weight values.



Show test data

Discretize output

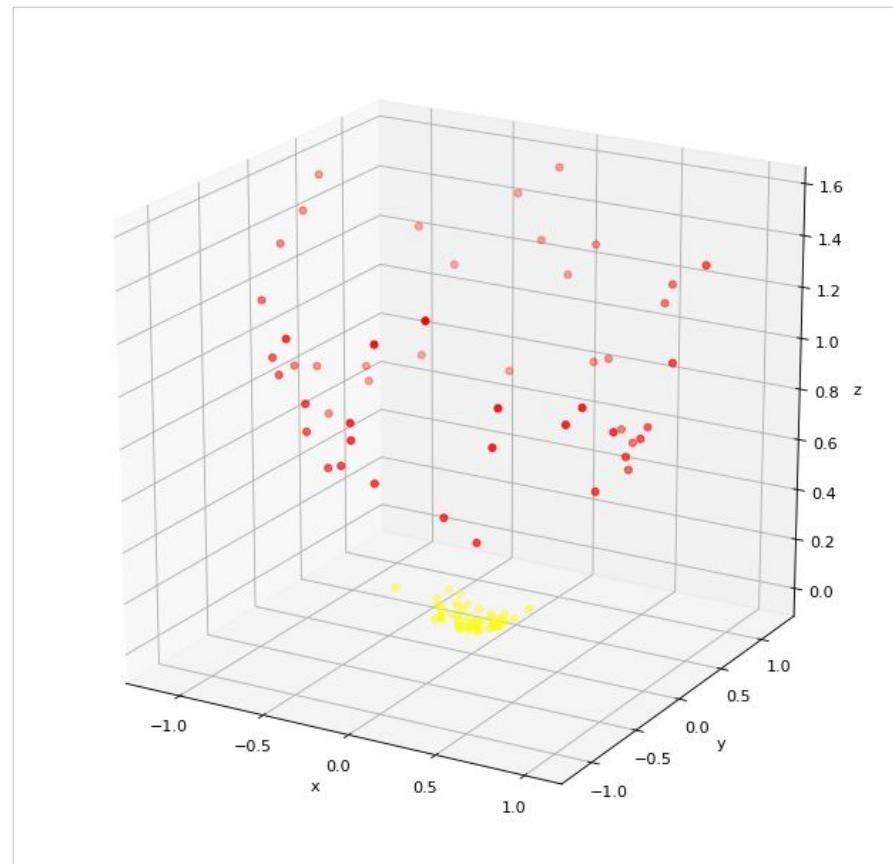
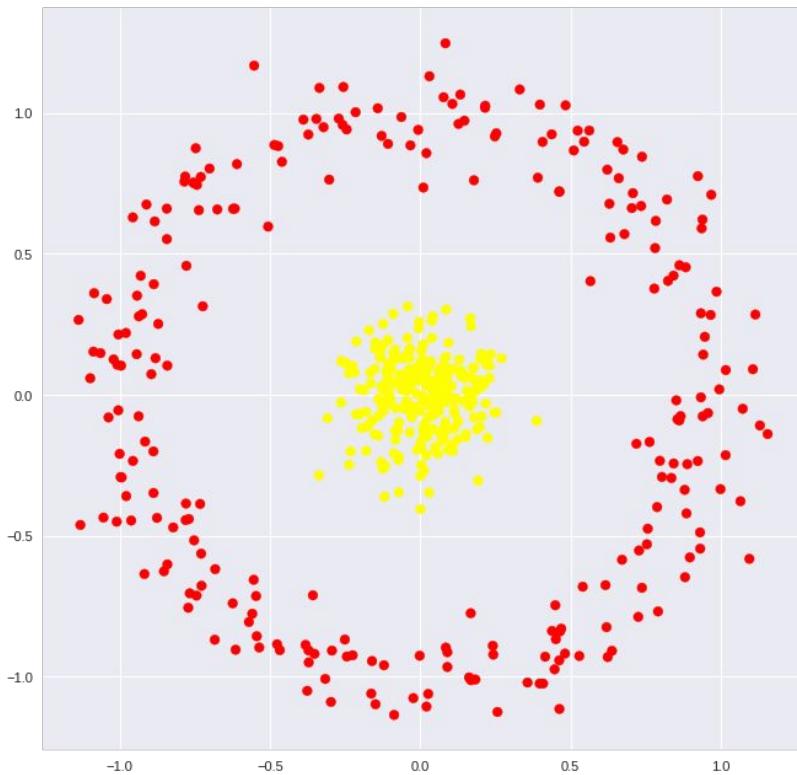


Feature engineering. What if we add a new feature:

$$z = x^2 + y^2.$$

Intuition

- All values of z will be positive.
- Yellow points are closer to the origin...
- So sum of their squared coords will be lower than red!



DATA

Which dataset do you want to use?



Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES

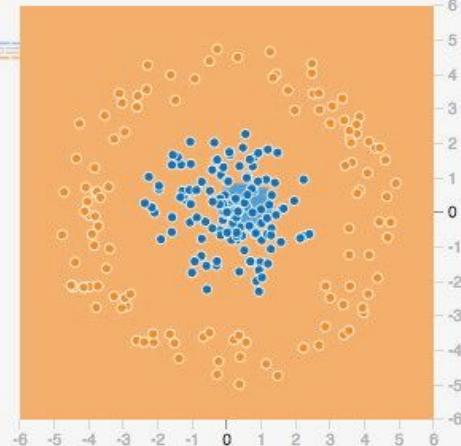
Which properties do you want to feed in?



0 HIDDEN LAYERS

OUTPUT

Test loss 0.486
Training loss 0.543



Colors shows
data, neuron and
weight values.



Show test data

Discretize output

DATA

Which dataset do you want to use?



Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?

X_1



X_2



X_1^2



X_2^2



$X_1 X_2$



$\sin(X_1)$



$\sin(X_2)$



+

-

1 HIDDEN LAYER

+

-

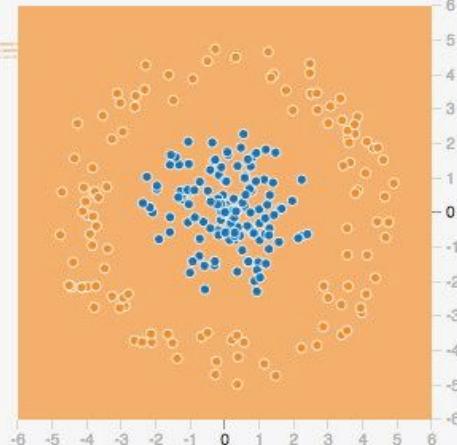
3 neurons

This is the output from one neuron.
Hover to see it larger.

OUTPUT

Test loss 0.547

Training loss 0.600



Colors shows
data, neuron and
weight values.



Show test data

Discretize output

DATA

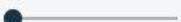
Which dataset do you want to use?



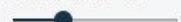
Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?



4 HIDDEN LAYERS

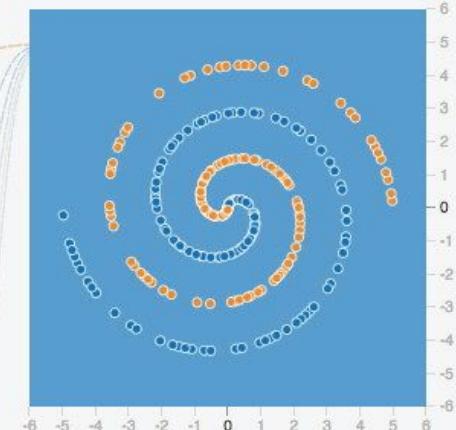


4 HIDDEN LAYERS

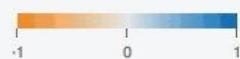
OUTPUT

Test loss 0.511

Training loss 0.502



Colors shows data, neuron and weight values.

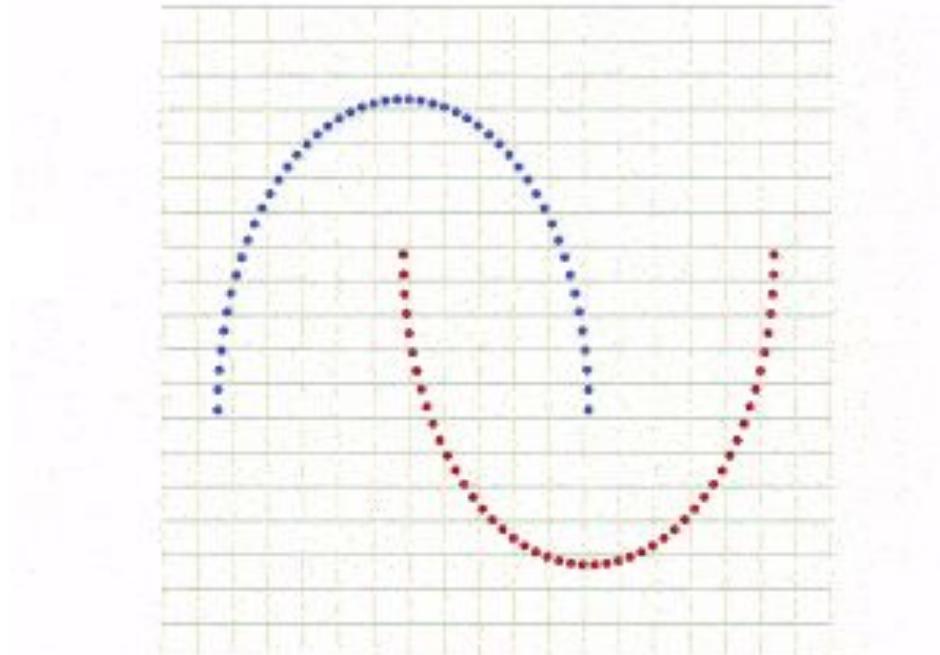


Show test data

Discretize output

This is the output from one **neuron**. Hover to see it larger.

The outputs are mixed with varying **weights**, shown by the thickness of the lines.



Exercise 1

Linear regression in TensorFlow 2 (low-level,
so you can see every piece involved).



Exercise 1

Goals

- Install TensorFlow 2.0
- Introduce Colaboratory
- Introduce the ingredients needed to train a neural network (training and testing data, **prediction**, **loss**, and **optimization**)

bit.ly/tf-ws1



For beginners and experts

For beginners

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

TF 1.x

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

TF 2.0

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

TR



K



Keras and tf.keras

The **clearest Deep Learning library** that exists today.

- For fast prototyping, advanced research, and production.

keras.io = reference implementation

- `import keras`

tf.keras = TensorFlow's implementation (a superset, built-in to TF, no need to install Keras separately)

- `from tensorflow import keras`

Exercise 2

Fashion MNIST in TensorFlow 2.0



Exercise 2

Goals

- Learn about the Sequential API
- Train a simple image classifier

Visit

bit.ly/tf-ws4

For experts

```
class MyModel(tf.keras.Model):

    def __init__(self, num_classes=10):
        super(MyModel, self).__init__(name='my_model')
        self.dense_1 = layers.Dense(32, activation='relu')
        self.dense_2 = layers.Dense(num_classes, activation='sigmoid')

    def call(self, inputs):
        # Define your forward pass here,
        x = self.dense_1(inputs)
        return self.dense_2(x)
```



What's the difference?



Symbolic vs Imperative APIs

Symbolic (Keras Sequential)

- Your model is a graph of layers
- Any graph you compile will run
- **TensorFlow helps you debug** by catching errors at **compile time**



Symbolic vs Imperative APIs

Symbolic (Keras Sequential)

- Your model is a graph of layers
- Any graph you compile will run
- **TensorFlow helps you debug** by catching errors at **compile time**

Imperative (Keras Subclassing)

- Your model is Python bytecode
- Complete flexibility and control
- Harder to debug / **harder to maintain**

Use a built-in training loop...

```
model.fit(x_train, y_train, epochs=5)
```

Or define your own

```
model = MyModel()

with tf.GradientTape() as tape:
    logits = model(images)
    loss_value = loss(logits, labels)

grads = tape.gradient(loss_value, model.trainable_variables)
optimizer.apply_gradients(zip(grads, model.trainable_variables))
```

TensorBoard

```
tb_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir)
```

```
model.fit(  
    x_train, y_train, epochs=5,  
    validation_data=[x_test, y_test],  
    callbacks=[tb_callback])
```

- Show data download links
 Ignore outliers in chart scaling

Tooltip sorting method: default ▾

Smoothing



Horizontal Axis

STEP RELATIVE WALL

Runs

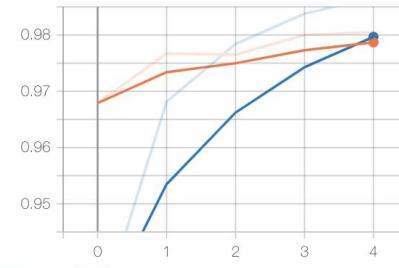
Write a regex to filter runs

- 20190227-033014/test
 20190227-033014/train

Filter tags (regular expressions supported)

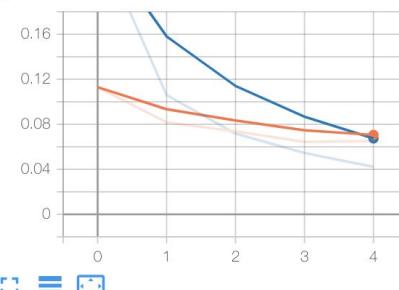
accuracy

accuracy
tag: accuracy



loss

loss
tag: loss



sequential

dense_2

dropout_1

dense_1

dropout

dense

flatten

Sequential Model Diagram

Exercise 3

Keras model subclassing and TensorBoard



Exercise 3

Goals

- Learn about the Subclassing API
- Run TensorBoard in the browser

bit.ly/tf-ws3

Code corrections

```
!pip install -q tensorflow==2.0.0-alpha0
```

```
!pip install -q tensorboard==1.13.0
```

```
%load_ext tensorboard.notebook (thank you!)
```



TensorFlow Extended (TFX)

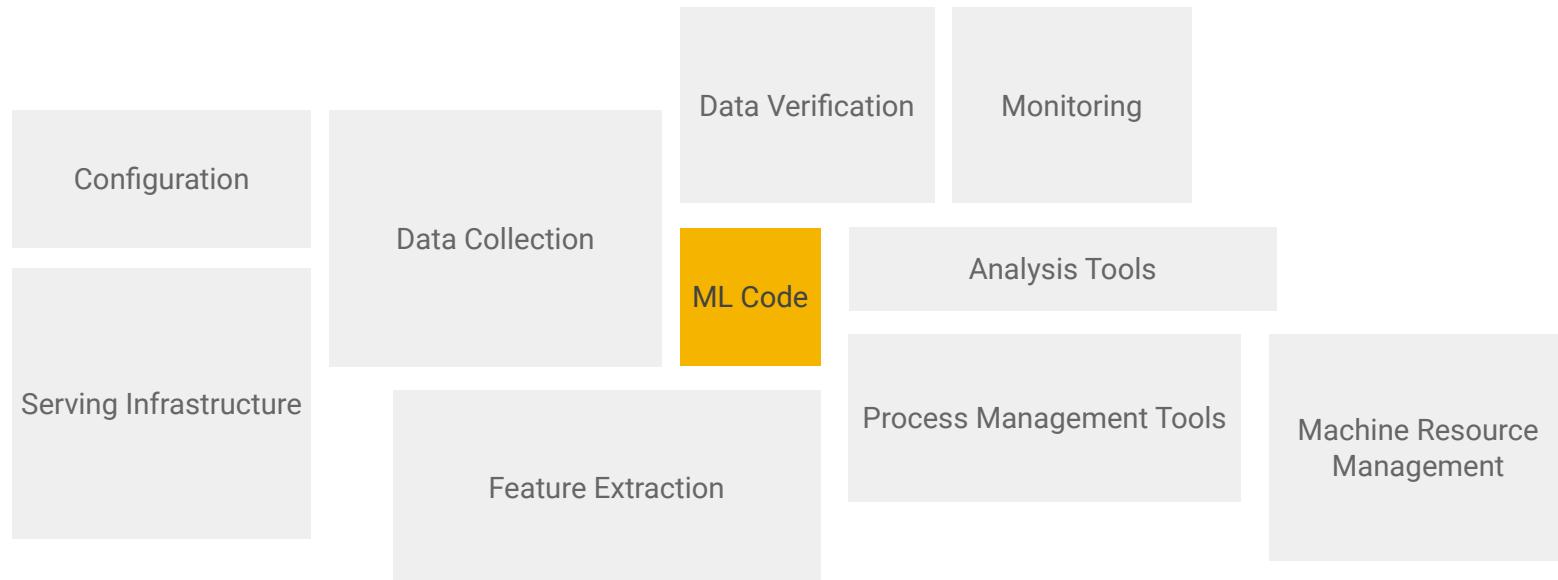
Real World Machine Learning in Production

In addition to training a model ...

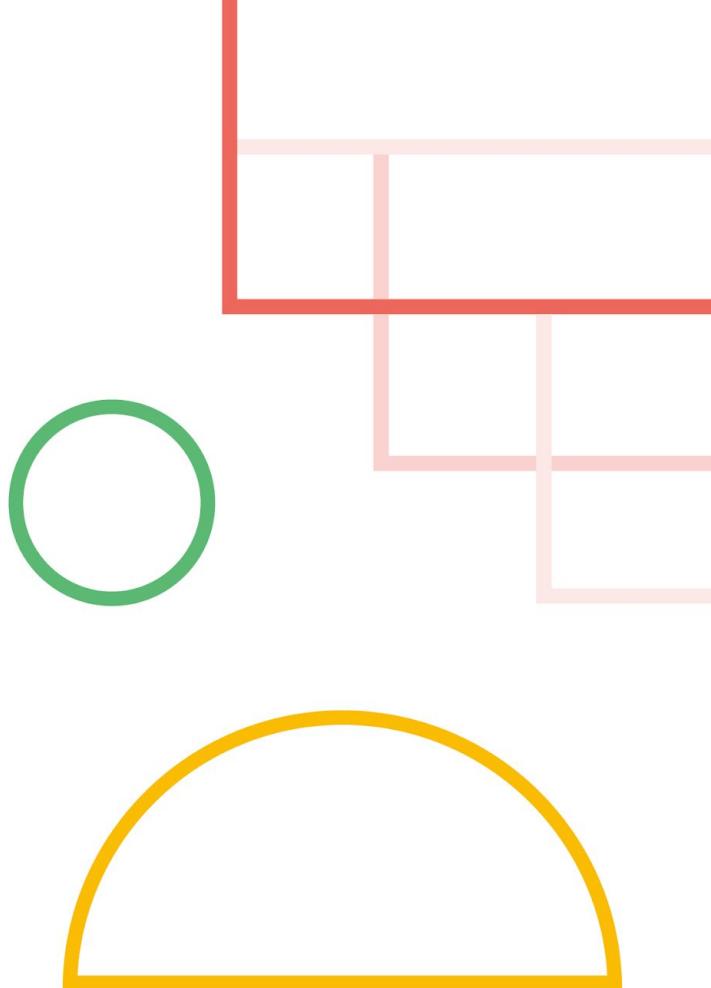
A solid yellow square containing the text "ML Code".

ML Code

... a production solution requires so much more



Tensorflow Extended (TFX)



TensorFlow Extended (TFX)

Powers our most important bets and products





And some of our most
important partners





Production Machine Learning

Machine Learning Development

- Labeled data
- Feature space coverage
- Minimal dimensionality
- Maximum predictive data
- Fairness
- Rare conditions
- Data lifecycle management



Production Machine Learning

Machine Learning Development

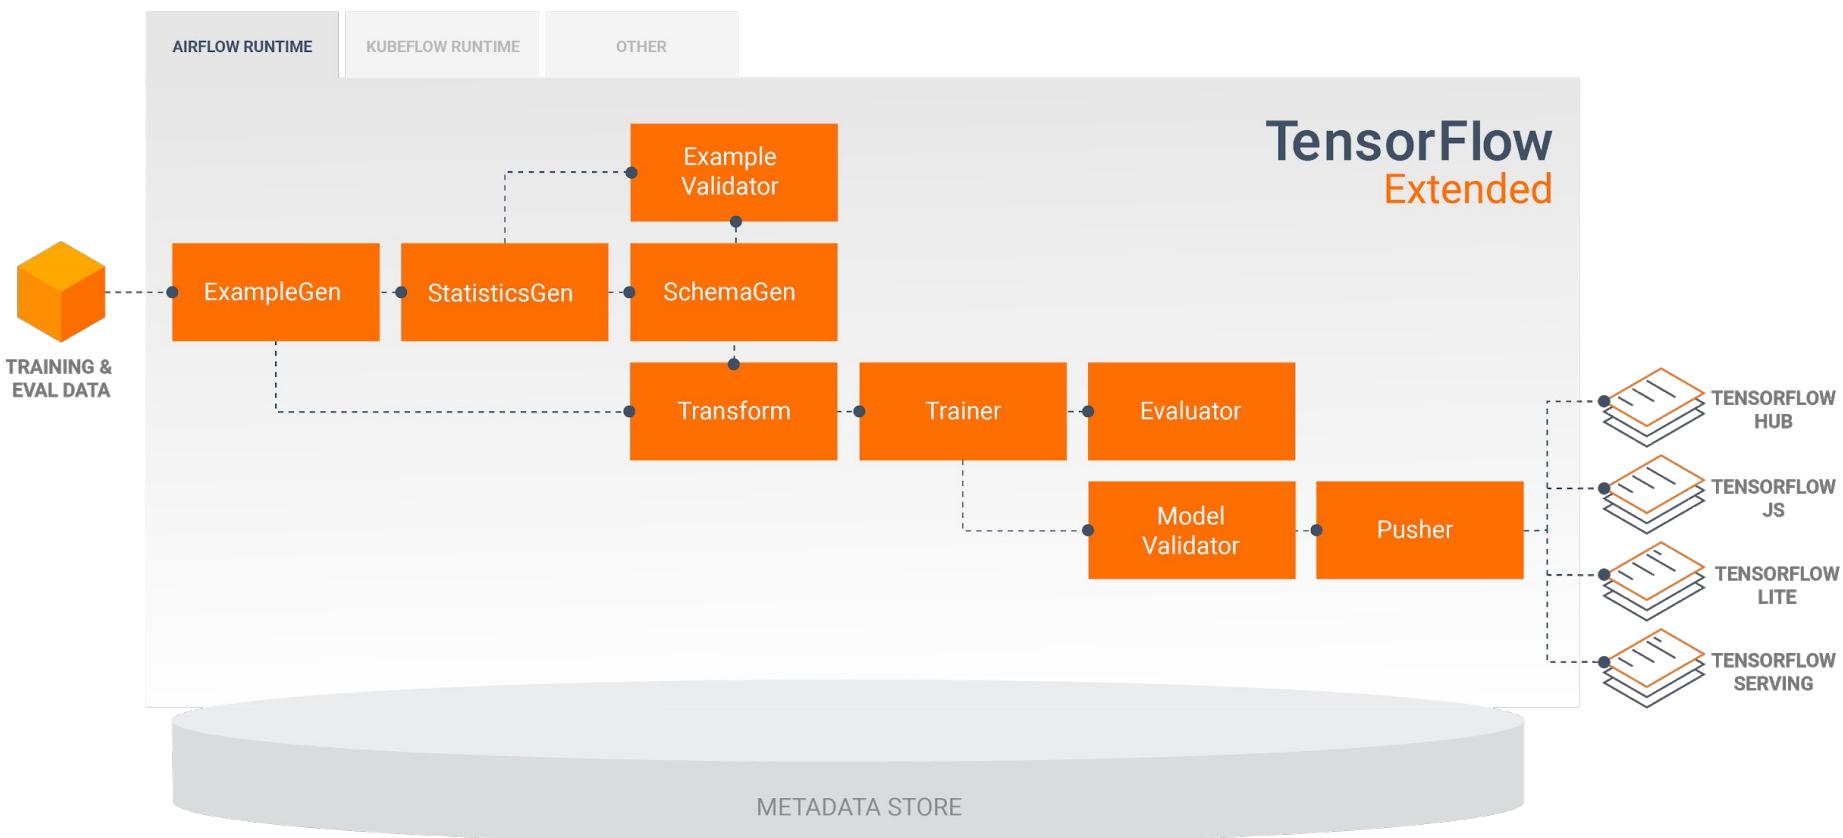
- Labeled data
- Feature space coverage
- Minimal dimensionality
- Maximum predictive data
- Fairness
- Rare conditions
- Data lifecycle management



Modern Software Development

- Scalability
- Extensibility
- Configuration
- Consistency & Reproducibility
- Modularity
- Best Practices
- Testability
- Monitoring
- Safety & Security

TFX CONFIG





TensorFlow Serving

Running your trained model in a server/cluster



Key Concepts

- **Servables** are the central abstraction in TensorFlow Serving. Servables are the underlying objects that clients use to perform computation (for example, a lookup or inference)



Key Concepts

Servable Versions

TensorFlow Serving can handle one or more versions of a servable over the lifetime of a single server instance.

Servable Streams

A servable stream is the sequence of versions of a servable, sorted by increasing version numbers.



Key Concepts

Models

TensorFlow Serving represents a model as one or more servables.

Loaders

Loaders manage a servable's life cycle. The Loader API enables common infrastructure independent from specific servables.



Key Concepts

Sources

Sources are plugin modules that find and provide servables. Each Source provides zero or more servable streams. For each servable stream, a Source supplies one Loader instance for each version it makes available to be loaded.



Key Concepts

Aspired Versions

Aspired versions represent the set of servable versions that should be loaded and ready.



Key Concepts

Version Policy

Version Policies specify the sequence of version loading and unloading within a single servable stream.

- **Availability Preserving Policy** - *Avoid leaving zero versions loaded; typically load a new version before unloading an old one*
- **Resource Preserving Policy** - *Avoid having two versions loaded simultaneously, thus requiring double the resources (unload an old version before loading a new one)*
- **Custom** - Whatever you need



Key Concepts

Managers

Managers handle the full lifecycle of Servables, including:

- loading Servables
- serving Servables
- unloading Servables



Key Concepts

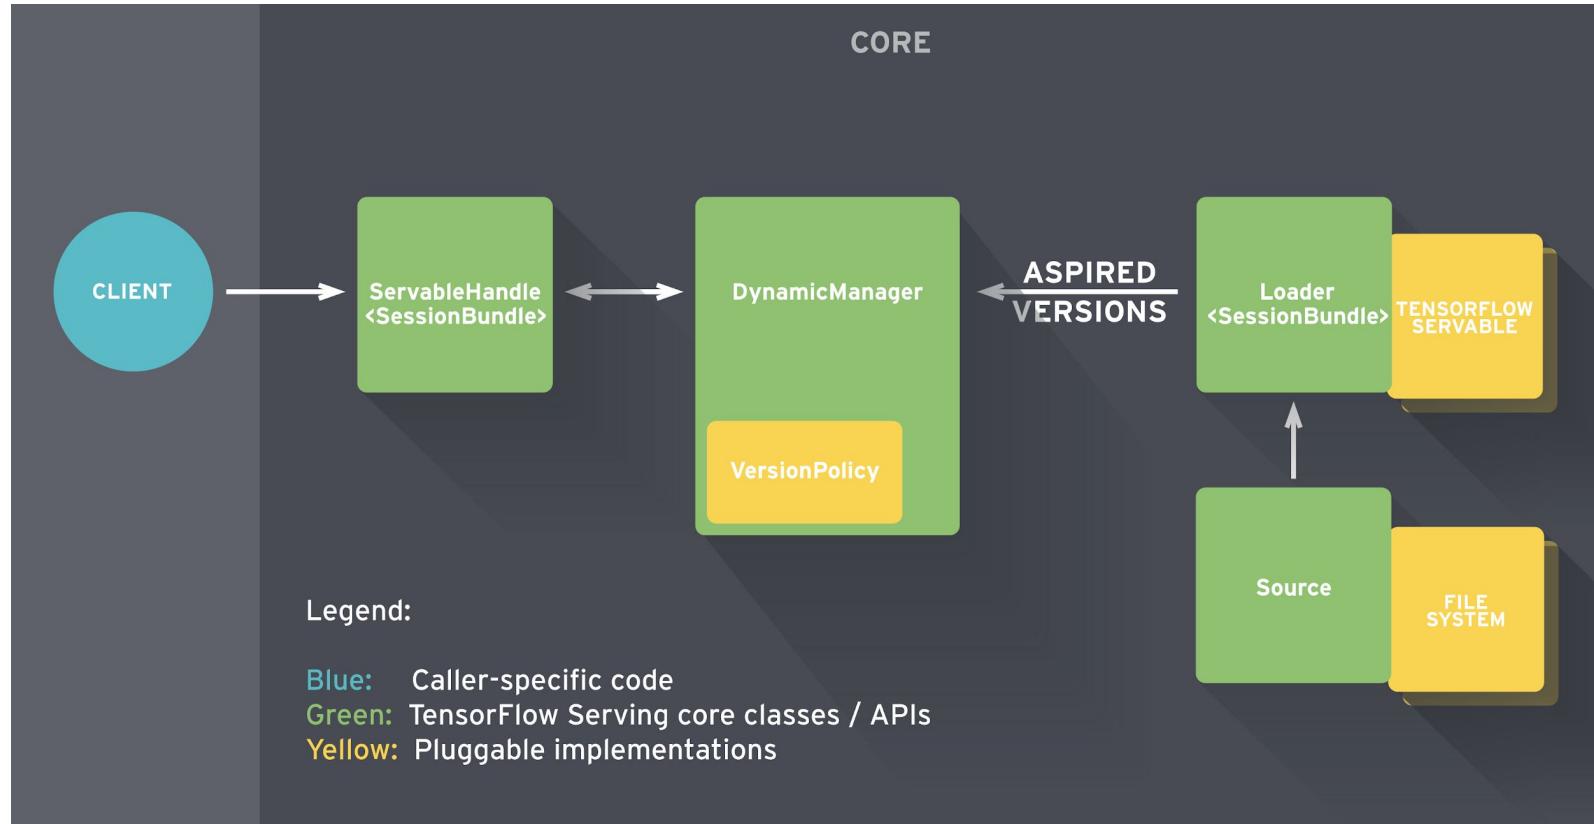
ServableHandle

Managers provide a simple, narrow interface -- `GetServableHandle()` -- for clients to access loaded servable instances.

Clients ask the Manager for the Servable, either specifying a version explicitly or just requesting the latest version. The Manager returns a handle for the Servable.



Architecture





Exercise

- Train a model
- Save a SavedModel
- Serve it

<http://bit.ly/serving-colab>

Beyond Hello World





A few of my favorites

- Machine Translation
- Image Captioning (the decoder is similar!)
- DCGan and Pix2Pix



The docs are runnable

Tutorials on tf.org/alpha are

- Backed by a Jupyter Notebook
- Can be run directly in Colab

They automatically

- Install the right TensorFlow version
- Download a dataset
- Train a model
- Show you the result

tensorflow.org/alpha/tutorials/text/image_captioning

TensorFlow > Learn > TensorFlow Core > TF 2.0 Alpha

Image Captioning with Attention

Run in Google Colab

View source on GitHub

Given an image like the below, our goal is to generate a caption, such as "a surfer riding on a wave".





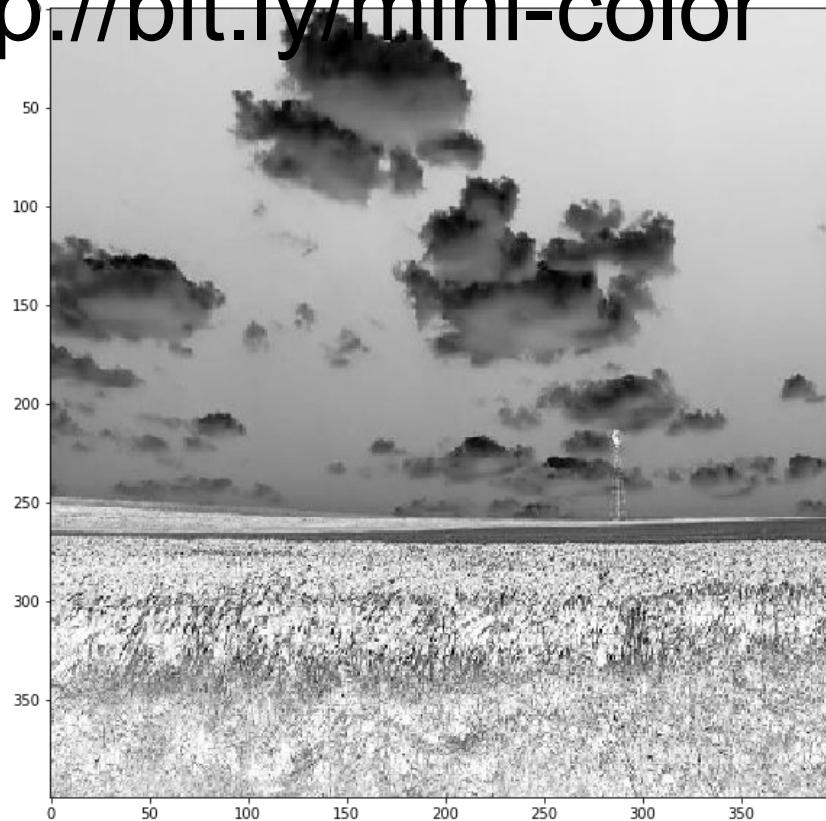
<http://bit.ly/mini-dream>



<https://github.com/random-forests/applied-dl/blob/master/examples/9-deep-dream-minimal.ipynb>

Code walkthrough

 <http://bit.ly/mini-color>



<https://github.com/random-forests/applied-dl/blob/master/examples/9-image-colorization.ipynb>

Exercise 6

Deep Dream



Is anyone bilingual? Trilingual?

When translating, do you...

- Go directly from source -> target
- Or, go from source -> **intermediate representation** -> target.



Machine translation tutorials

- [Hello world](#) (seq2seq), trains in about a minute.
- [Neural Machine Translation with Attention](#)
- [Transformer](#)

P.S., isn't 2019 cool? It's **amazing** this is possible.



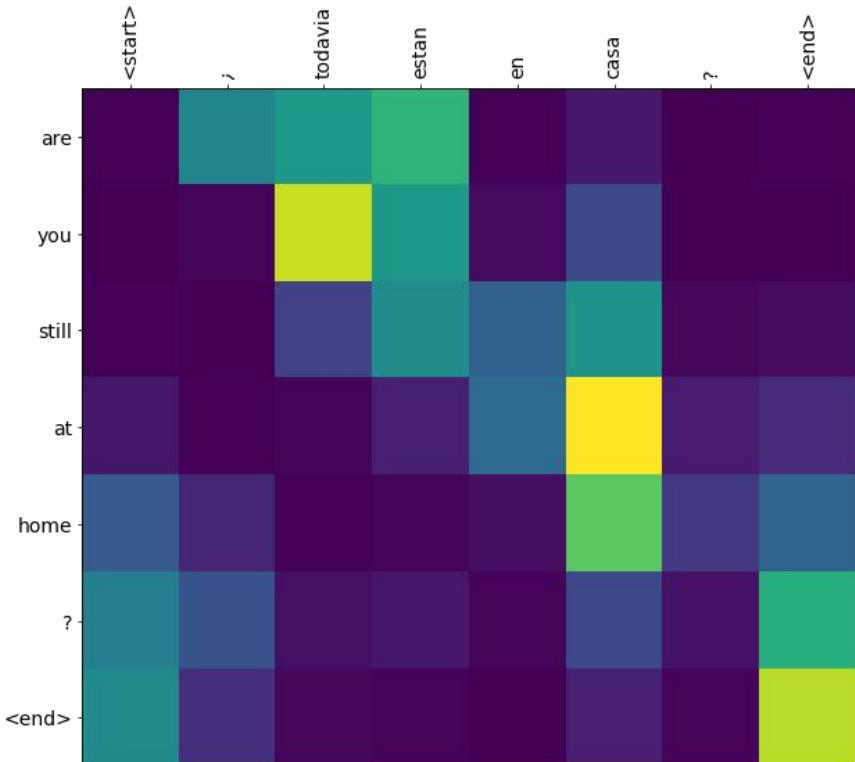
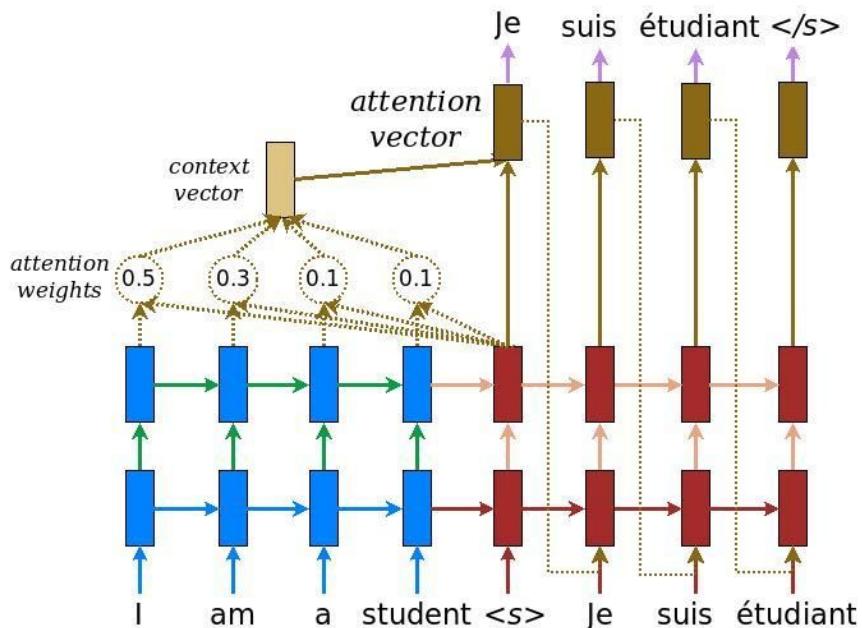
Exercise 7

Goals

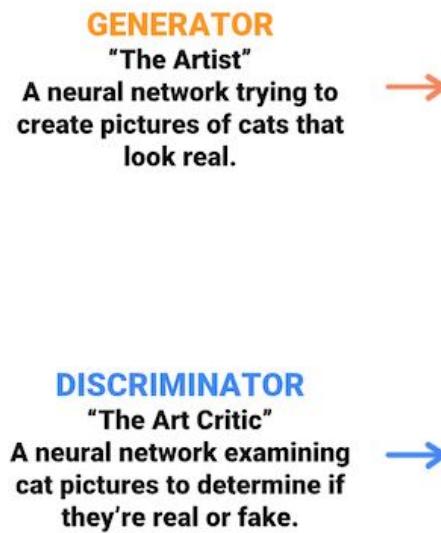
- Train an English to Spanish model, just for fun
- Learn about encoder / decoders

Visit

bit.ly/minimal-nmt

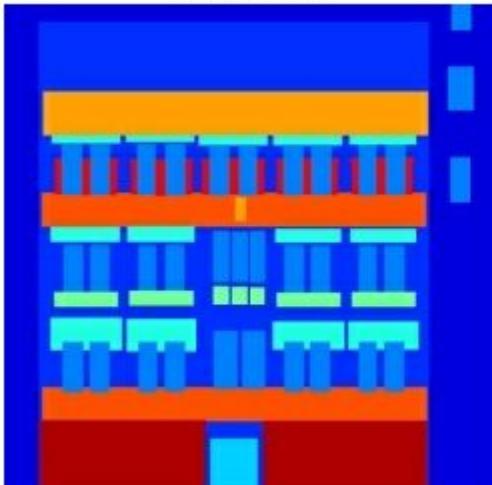


https://www.tensorflow.org/alpha/tutorials/sequences/nmt_with_attention



<https://www.tensorflow.org/alpha/tutorials/generative/dcgan>

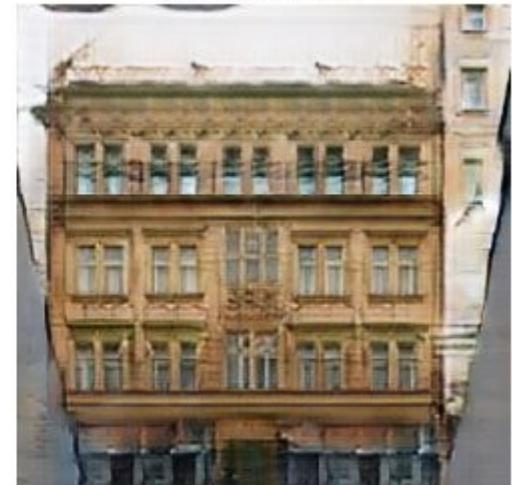
Input Image



Ground Truth



Predicted Image



<https://www.tensorflow.org/alpha/tutorials/generative/pix2pix>



Prediction Caption: the person is riding a surfboard in the ocean <end>

https://www.tensorflow.org/alpha/tutorials/sequences/image_captioning



Exercise 6

Goals

- Use a pretrained CNN
- Extract intermediate activations
- Compute gradients w.r.t. an image

Visit

bit.ly/tf-ws6

Under the hood



Let's make this faster

```
lstm_cell = tf.keras.layers.LSTMCell(10)

def fn(input, state):
    return lstm_cell(input, state)

input = tf.zeros([10, 10]); state = [tf.zeros([10, 10])] * 2
lstm_cell(input, state); fn(input, state) # warm up

# benchmark
timeit.timeit(lambda: lstm_cell(input, state), number=10) # 0.03
```

Let's make this faster

```
lstm_cell = tf.keras.layers.LSTMCell(10)

@tf.function
def fn(input, state):
    return lstm_cell(input, state)

input = tf.zeros([10, 10]); state = [tf.zeros([10, 10])] * 2
lstm_cell(input, state); fn(input, state) # warm up

# benchmark
timeit.timeit(lambda: lstm_cell(input, state), number=10) # 0.03
timeit.timeit(lambda: fn(input, state), number=10) # 0.004
```

AutoGraph makes this possible

```
@tf.function
def f(x):
    while tf.reduce_sum(x) > 1:
        x = tf.tanh(x)
    return x

# you never need to run this (unless curious)
print(tf.autograph.to_code(f))
```

Generated code

```
def tf_f(x):
    def loop_test(x_1):
        with ag__.function_scope('loop_test'):
            return ag__.gt(tf.reduce_sum(x_1), 1)
    def loop_body(x_1):
        with ag__.function_scope('loop_body'):
            with ag__.utils.control_dependency_on_returns(tf.print(x_1)):
                tf_1, x = ag__.utils.alias_tensors(tf, x_1)
                x = tf_1.tanh(x)
            return x,
    x = ag__.while_stmt(loop_test, loop_body, (x,), (tf,))
    return x
```

Going big: tf.distribute.Strategy

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(64, input_shape=[10]),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')))

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

Going big: Multi-GPU

```
strategy = tf.distribute.MirroredStrategy()

with strategy.scope():

    model = tf.keras.models.Sequential([
        tf.keras.layers.Dense(64, input_shape=[10]),
        tf.keras.layers.Dense(64, activation='relu'),
        tf.keras.layers.Dense(10, activation='softmax')])

    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
```



What's different between TF1 and TF2?

Removed

- `session.run`
- `tf.control_dependencies`
- `tf.global_variables_initializer`
- `tf.cond`, `tf.while_loop`

Added

- `tf.function`, AutoGraph

TensorFlow.js



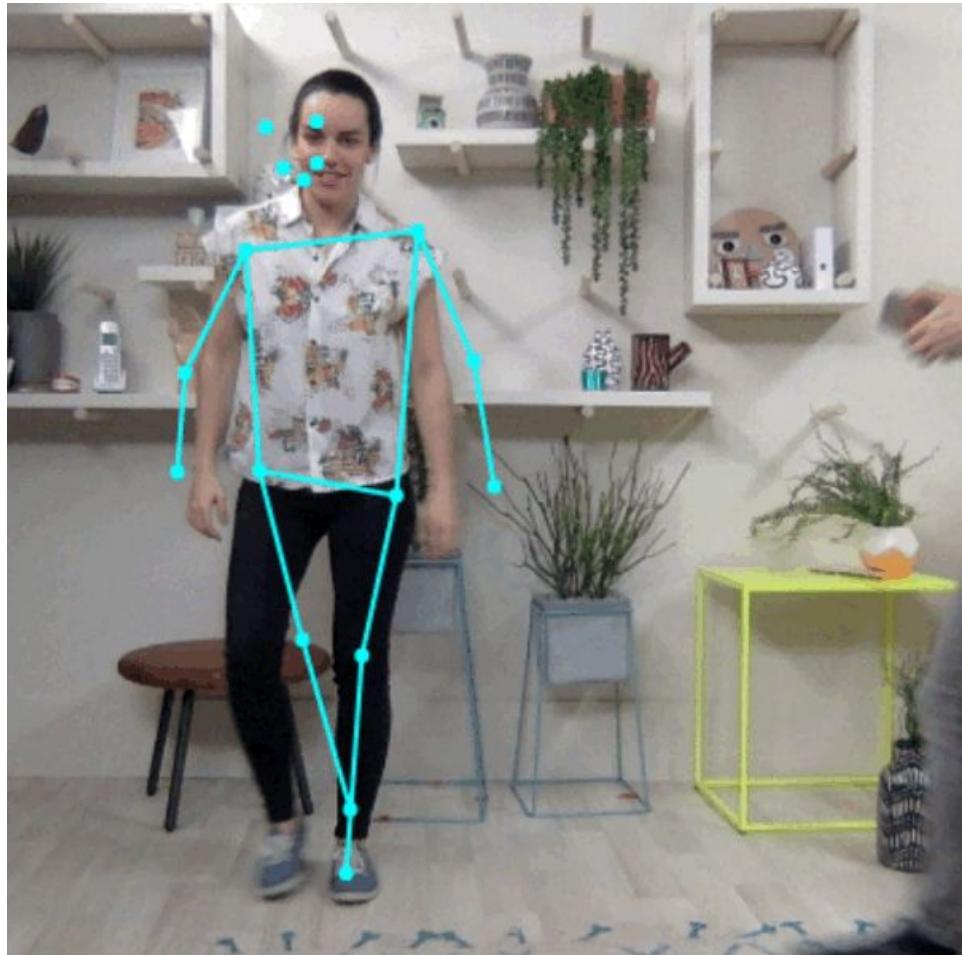
Demo #1

PoseNet



PoseNet

bit.ly/pose-net



Demo #2

BodyPix



BodyPix

bit.ly/body-pix



Exercise 10

Deploy a model in the browser



Exercise 10

Goals

- Train a model in Python and run it in a webpage, without leaving Colab

Visit

bit.ly/c-to-web

Learning more





Learn more

Tutorials and guides

- tensorflow.org/alpha

Books

- [Deep Learning with Python](#)
- Hands-On Machine Learning with Scikit-Learn and TensorFlow
(version 2.0 only, in +/- a month or two)

Courses

- [Intro to Deep Learning](#) (MIT)
- [Convolutional Neural Networks for Visual Recognition](#) (Stanford)



tf.thanks!

Josh Gordon (twitter.com/random_forests)

Robert Crowe (twitter.com/robert_crowe)

Extras



Exercise 4

Structured data



Exercise 4

Goals

- Start from raw data
- Create a model

Visit

bit.ly/tf-ws4a

Facets pair-code.github.io/facets/